

# **Session 4: MPI Parallel-Distributed Computation on Kubernetes**

**Sunday, August 5, 2018 · 2:30PM - 3:30PM**

**Christopher Paolini**

**Assistant Professor of Electrical and Computer Engineering**

**San Diego State University**

**PI NSF OAC CC\*Storage #1659169, Co-PI NSF OAC CC-NIE #1245312,**


**Co-PI NSF OAC MRI #0922702, Co-PI NSF OAC CI-TEAM #0753283**

# Topics

- **Configuring and testing an environment for using OpenMPI on K8s**
- **Benchmarking PRP K8s with LINPACK**
- **Example PRP science driver: subsurface CO<sub>2</sub> and waste-water injection simulation**

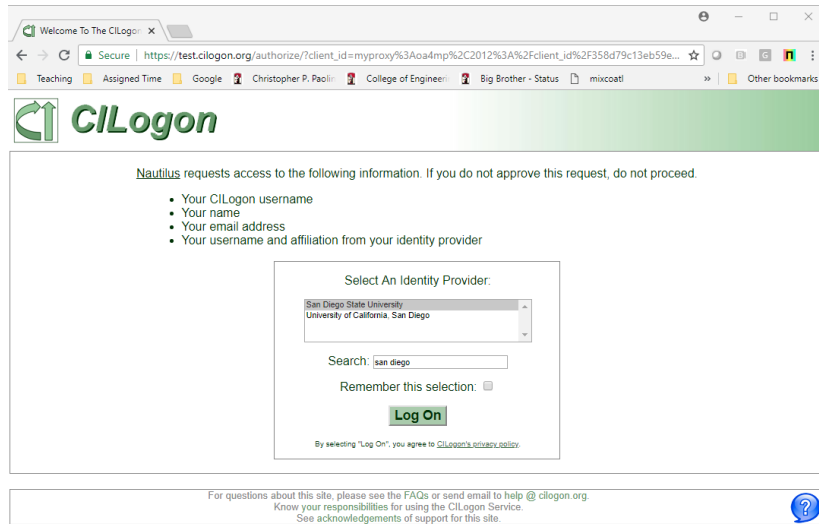
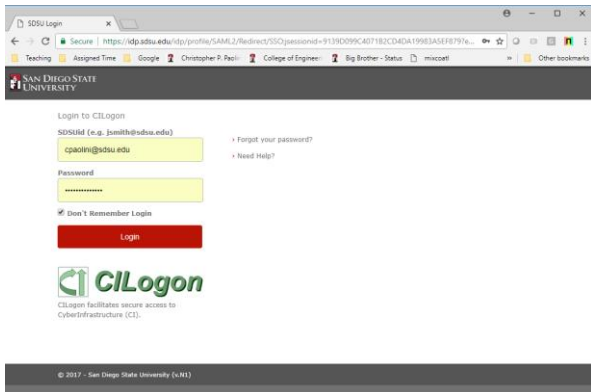
# PRP Kubernetes Gateway: *kubectl* Client Config

- Login to <https://nautilus.optiputer.net/> to obtain the latest *kubectl* client config bundle



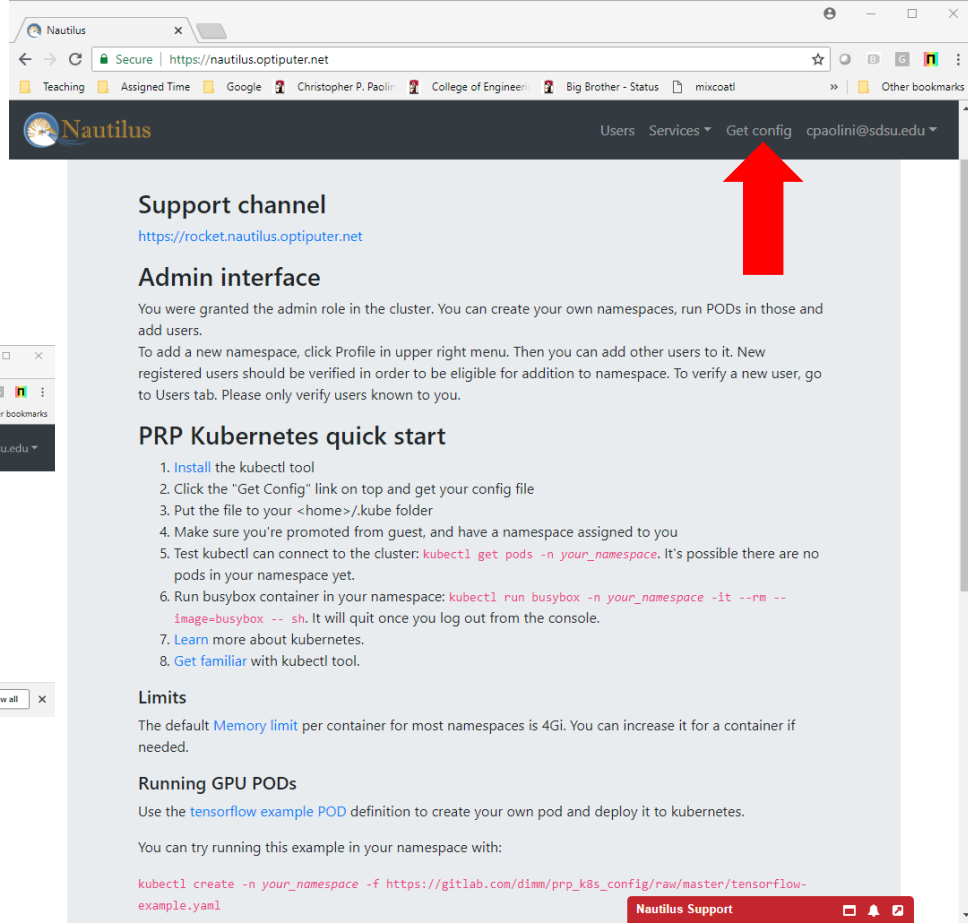
Welcome to the PRP Kubernetes gateway

Nautilus Support



# PRP Kubernetes Gateway: *kubectl* Client Config cont.

- Download **kubectl** client config bundle via URL <https://nautilus.optiputer.net/authConfig> and move file to `$HOME/.kube/config` on your station



The screenshot shows the Nautilus web interface in a browser. The address bar displays `https://nautilus.optiputer.net`. The navigation bar includes links for 'Users', 'Services', 'Get config', and the user 'cpalini@sdsu.edu'. A red arrow points to the 'Get config' link. The main content area includes a 'Support channel' link, an 'Admin interface' section with instructions on adding users, and a 'PRP Kubernetes quick start' section with an 8-step guide. A 'Limits' section mentions a 4Gi default memory limit. The 'Running GPU PODs' section provides a command to create a pod. At the bottom, there is a 'Nautilus Support' button.

Support channel  
<https://rocket.nautilus.optiputer.net>

Admin interface  
You were granted the admin role in the cluster. You can create your own namespaces, run PODs in those and add users.  
To add a new namespace, click Profile in upper right menu. Then you can add other users to it. New registered users should be verified in order to be eligible for addition to namespace. To verify a new user, go to Users tab. Please only verify users known to you.

PRP Kubernetes quick start


1. [Install](#) the **kubectl** tool
2. Click the "Get Config" link on top and get your config file
3. Put the file to your `<home>/.kube` folder
4. Make sure you're promoted from guest, and have a namespace assigned to you
5. Test **kubectl** can connect to the cluster: `kubectl get pods -n your_namespace`. It's possible there are no pods in your namespace yet.
6. Run **busybox** container in your namespace: `kubectl run busybox -n your_namespace -it --rm --image=busybox -- sh`. It will quit once you log out from the console.
7. [Learn](#) more about **kubernetes**.
8. [Get familiar](#) with **kubectl** tool.

Limits  
The default **Memory limit** per container for most namespaces is 4Gi. You can increase it for a container if needed.

Running GPU PODs  
Use the [tensorflow example POD](#) definition to create your own pod and deploy it to **kubernetes**.  
You can try running this example in your namespace with:  

```
kubectl create -n your_namespace -f https://gitlab.com/dimm/prp_k8s_config/raw/master/tensorflow-example.yaml
```

Nautilus Support



Please put the downloaded 'config' file into your `~/kube/` folder

# Using OpenMPI on K8s: *Create Pods*

- Check the kubectl configuration

```
[paolini@ps-40g ~]$ kubectl -n sdsu cluster-info
```

Kubernetes master is running at https://67.58.53.147:6443

- Create cluster of pods using a YAML configuration

```
[paolini@fiona k8s]$ kubectl create -f subflow.yaml
```

- Verify all pods are in the *Running* state

```
[paolini@fiona k8s]$ kubectl get pods -o wide -n sdsu
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
subflow-74c57d67d4-bqnc2	1/1	Running	0	3d	10.244.1.151	k8s-nvme-01.sdsc.optiputer.net
subflow-74c57d67d4-gd8sn	1/1	Running	0	3d	10.244.16.45	fiona.cac.washington.edu
subflow-74c57d67d4-jjxl5	1/1	Running	0	3d	10.244.15.182	dtn-main.ucr.edu
subflow-74c57d67d4-jl72q	1/1	Running	0	3d	10.244.11.71	siderea.ucsc.edu
subflow-74c57d67d4-l9n8t	1/1	Running	0	3d	10.244.24.46	dtn2-daejeon.kreonet.net
subflow-74c57d67d4-lj8xw	1/1	Running	0	3d	10.244.10.96	fiona-dtn-1.ucsc.edu
subflow-74c57d67d4-q8mbj	1/1	Running	0	3d	10.244.19.134	fiona.nwsc.ucar.edu
subflow-74c57d67d4-x5j6r	1/1	Running	0	3d	10.244.12.100	k8s-epyc-01.sdsc.optiputer.net

# Using OpenMPI on K8s: *Create Pods cont.*

```
[paolini@fiona k8s]$ kubectl get pod -n sdsu -o=custom-columns=NODE:.spec.nodeName
```

NODE

k8s-nvme-01.sdsc.optiputer.net

fiona.cac.washington.edu

dtm-main.ucr.edu

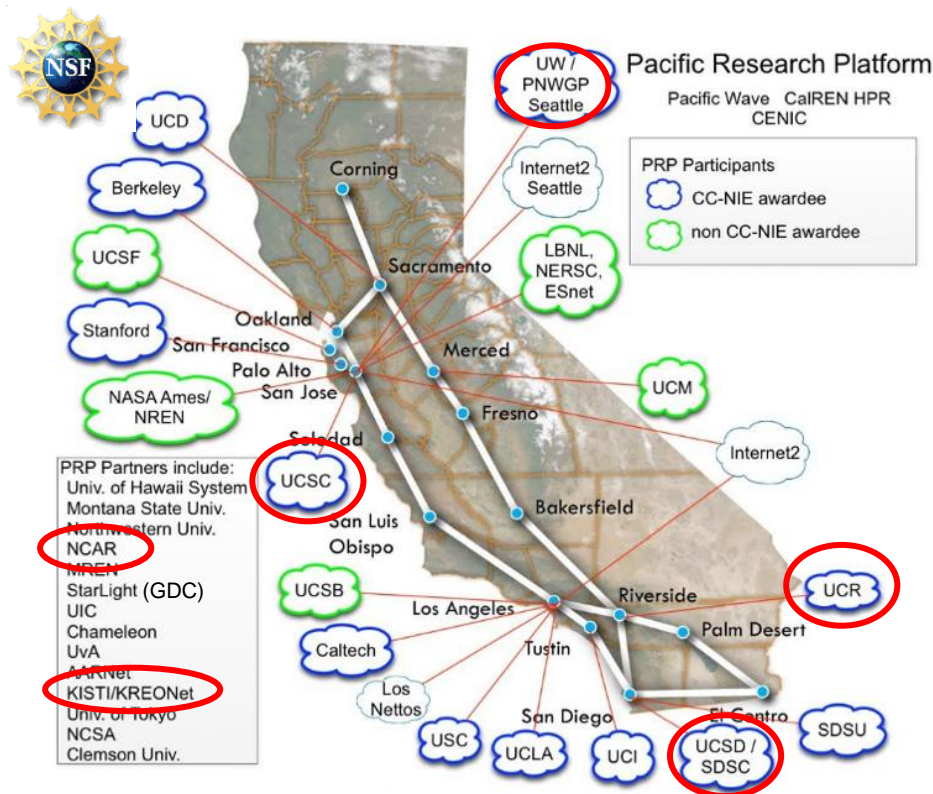
siderea.ucsc.edu

dtm2-daejeon.kreonet.net

fiona-dtm-1.ucsc.edu

fiona.nwsc.ucar.edu

k8s-epyc-01.sdsc.optiputer.net



Note: this diagram represents a subset of sites and connections.

v1.16 - 20151019

# Using OpenMPI on K8s: *YAML configuration*

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: subflow
  namespace: sdsu
spec:
  replicas: 8
  template:
    metadata:
      labels:
        k8s-app: subflow
    spec:
      containers:
        - name: subflow
          image: phusion/baseimage:0.9.19
          imagePullPolicy: IfNotPresent
          args: ["sleep", "infinity"]
          resources:
            limits:
              memory: "48Gi"
            requests:
              memory: "32Gi"
          volumeMounts:
            - name: nfs
              mountPath: /nfs
      volumes:
        - name: nfs
          nfs:
            server: 10.109.158.238
            path: "/"
```

Request each container to have 32GiB of allocatable memory, with an upper limit of 48GiB

Mount existing NFS (Network File System) share (nfs://10.109.158.238/) on mount-point /nfs in pod.

Persistent location of git source clone and build for k8s pods

# Using OpenMPI on K8s: *Initialize Cluster*

```
[paolini@fiona k8s]$ ./initcluster.sh
[paolini@fiona k8s]$ cat initcluster.sh
#!/bin/bash
hosts=""
n=0
for host in `kubectl get pod -n sdsu -o=custom-columns=IP:status.podIP,NAME:.metadata.name | grep subflow`
do
    hosts="$hosts $host"
    if [ $n -eq 1 ];
    then
        cores=`kubectl exec -it $host -n sdsu -- bash -c "grep -c ^processor /proc/cpuinfo"`
        cores=`tr -dc '[:print:]' <<< " $cores"`
        hosts+=$cores
        n=0
    else
        n=1
    fi
done
pods=`kubectl get pod -n sdsu -o=custom-columns=NAME:.metadata.name|sed -e '/NAME/d'|egrep '^subflow-'`
for pod in $pods
do
    kubectl exec -it $pod -n sdsu -- bash -c "/nfs/subflow/k8s/setup.sh $hosts"
done
for pod in $pods
do
    kubectl exec -it $pod -n sdsu -- bash -c "/nfs/subflow/k8s/updatekeys.sh"
done
```

**Build string array of 3-tuples (IP Address, hostname, processor count) needed for MPI *hostfile***

**Invoke setup.sh on each pod with string array as argument list**

**setup.sh starts ssh server on each pod**

**Gather ssh public keys of all pods for generating /etc/ssh/ssh\_known\_hosts (requires sshd be running on all pods)**



# Using OpenMPI on K8s: *Initialize Pod*

```
[paolini@fiona k8s]$ cat setup.sh
#!/bin/bash

apt-get --quiet update;
apt-get --quiet -y install screen emacs git python mlocate openmpi-bin openmpi-doc libopenmpi-dev make
valgrind g++ m4 gfortran liblapacke-dev libnetcdf-dev iputils-ping openssh-server cmake mesa-utils-extra

if [ ! -f /etc/hosts.DIST ]
then
    cp /etc/hosts /etc/hosts.DIST
fi

if ! id subflow >/dev/null 2>&1
then
    groupadd subflow; useradd -g subflow -d /home/subflow -m -c 'Subflow Execution' -s /bin/bash subflow
fi

if [ ! -f /root/.ssh/id_rsa.pub ]
then
    echo -e 'y\n'|/usr/bin/ssh-keygen -f /root/.ssh/id_rsa -t rsa -N ''
fi
```

**Install required packages for running MPI programs and building target**

**Backup original /etc/hosts file**

**Create fictitious user account if you don't want to run MPI processes as root**

**Generate protocol version 2 ssh authentication keys for passwordless, public key authentication between pods**

# Using OpenMPI on K8s: *Initialize Pod cont.*

```
c=`grep -c "HostbasedAuthentication yes" /etc/ssh/ssh_config`  
if [ "$c" -eq "0" ]  
then  
    echo "Enable SSH client HostbasedAuthentication"  
    echo "HostbasedAuthentication yes" >> /etc/ssh/ssh_config  
fi
```

```
c=`grep -c "HostbasedAuthentication yes" /etc/ssh/sshd_config`  
if [ "$c" -eq "0" ]  
then  
    echo "Enable SSH server HostbasedAuthentication"  
    echo "HostbasedAuthentication yes" >> /etc/ssh/sshd_config  
fi
```

```
c=`grep -c "IgnoreRhosts no" /etc/ssh/sshd_config`  
if [ "$c" -eq "0" ]  
then  
    echo "Enable SSH server rhosts authentication"  
    echo "IgnoreRhosts no" >> /etc/ssh/sshd_config  
fi
```

```
c=`grep -c "HostbasedUsesNameFromPacketOnly yes" /etc/ssh/sshd_config`  
if [ "$c" -eq "0" ]  
then  
    echo "Enable SSH server HostbasedUsesNameFromPacketOnly"  
    echo "HostbasedUsesNameFromPacketOnly yes" >> /etc/ssh/sshd_config  
fi
```

**Permit user authentication between pods listed in `/etc/ssh/shosts.equiv` (enable in client and server config files)**

**Enable authentication for root user, based on names of (trusted) pods in `/root/.rhosts`**

**Configure sshd to accept the hostname information provided in the connection itself, rather than use DNS resolution**

# Using OpenMPI on K8s: *Initialize Pod cont.*

```
c=`grep -c "^PermitRootLogin yes" /etc/ssh/sshd_config`  
if [ "$c" -eq "0" ]  
then  
    echo "Enable SSH server root login"  
    echo "PermitRootLogin yes" >> /etc/ssh/sshd_config  
fi
```

**Permit root authentication between pods to allow MPI processes to run as root**

```
cp /etc/hosts.DIST /etc/hosts  
rm -f /etc/ssh/shosts.equiv /etc/mpihostfile
```

```
while [[ $# -gt 0 ]]
```

```
do
```

```
    echo "$1 $2" >> /etc/hosts  
    echo "$2" >> /etc/ssh/shosts.equiv  
    echo "$2 slots=$3 max-slots=$3" >> /etc/mpihostfile  
    shift  
    shift  
    shift
```

```
done
```

```
cp /etc/ssh/shosts.equiv /root/.rhosts  
/etc/init.d/ssh stop  
/etc/init.d/ssh start
```

**Rebuild /etc/hosts, /etc/ssh/shosts.equiv, and /etc/mpihostfile using string array argument (IP Address, hostname, processor count)**

**Restart ssh server on pod to enable updated configuration**

# Using OpenMPI on K8s: *Initialize Pod cont.*

```
[paolini@fiona k8s]$ cat updatekeys.sh
```

```
#!/bin/bash
```

```
rm -f /etc/ssh/ssh_known_hosts
```

```
for host in `cat /etc/ssh/shosts.equiv`
```

```
do
```

```
    ssh-keyscan -t rsa $host >> /etc/ssh/ssh_known_hosts
```

```
    cp /etc/ssh/ssh_known_hosts /root/.ssh/known_hosts
```

```
done
```

Each host gathers the  
public ssh host keys of  
all hosts

# Using OpenMPI on K8s: *Using a Custom Image*

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: subflow
  namespace: sdsu
spec:
  replicas: 8
  template:
    metadata:
      labels:
        k8s-app: subflow
    spec:
      containers:
      - name: subflow
        image: dimm0/subflow:latest
        imagePullPolicy: IfNotPresent
        args: ["sleep", "infinity"]
        resources:
          limits:
            memory: "48Gi"
          requests:
            memory: "32Gi"
        volumeMounts:
        - name: nfs
          mountPath: /nfs
      volumes:
      - name: nfs
        nfs:
          server: 10.109.158.238
          path: "/"
```

```
[paolini@fiona image]$ cat Dockerfile
FROM phusion/baseimage:0.9.19


MAINTAINER Dmitry Mishin <dmishin@sdsc.edu>

RUN apt-get -y update && \
    apt-get -y install screen emacs git python mlocate openmpi-bin
openmpi-doc libopenmpi-dev make valgrind g++ m4 gfortran liblapack-dev
libnetcdf-dev iputils-ping openssh-server cmake mesa-utils-extra

[root@fiona image]# docker pull dimm0/subflow
Using default tag: latest
latest: Pulling from dimm0/subflow
f069f1d21059: Pull complete
ecbeec5633cf: Pull complete
ea6f18256d63: Pull complete
54bde7b02897: Pull complete
a3ed95caeb02: Pull complete
ce9e695a6234: Pull complete
346026b9659b: Pull complete
d8a2ef8e4be5: Pull complete
Digest:
sha256:eb39c3151cd0ed776f6188949bc7f93a64db2eb04ab4391591c2670e0c821806
Status: Downloaded newer image for dimm0/subflow:latest
```

# Using OpenMPI on K8s: *Using a Custom Container*

[root@fiona image]# docker images



REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
rocketchat/rocket.chat	latest	e98d34f3e7ab	16 hours ago	1.38GB
dimmm0/subflow	latest	6aab3414d3be	44 hours ago	1.17GB
us.gcr.io/prp-k8s/nautilus-portal	latest	e5f31849537f	7 days ago	38.7MB
us.gcr.io/prp-k8s/nautilus-portal	<none>	1c868229707b	7 days ago	38.7MB
us.gcr.io/prp-k8s/perfsonar_testpoint_h	latest	2c3315ea9b6d	10 days ago	818MB
us.gcr.io/prp-k8s/perfsonar_testpoint_h	<none>	3f220a2c3190	10 days ago	818MB
rocketchat/rocket.chat	<none>	587a250b077b	12 days ago	1.37GB
us.gcr.io/prp-k8s/meshconfig	latest	ee2c46848fde	2 weeks ago	34.2MB
sameersbn/postgresql	latest	3c0142eb3992	2 weeks ago	204MB
robcurrie/tensorflow-gpu	latest	90592bb75ee1	4 weeks ago	4.28GB
nvidia/k8s-device-plugin	1.10	a7c090961376	5 weeks ago	63.1MB
gcr.io/google_containers/kube-proxy-amd64	v1.11.0	1d3d7afd77d1	5 weeks ago	97.8MB
us.gcr.io/prp-k8s/prp-tuner	latest	62dbae4fe2e1	2 months ago	52.2MB
gcr.io/runconduit/proxy	v0.4.1	fd47f6c1933d	3 months ago	111MB
gcr.io/runconduit/proxy-init	v0.4.1	60eb863cee15	3 months ago	105MB
quay.io/calico/node	v3.1.1	d94b64ac210d	3 months ago	248MB
quay.io/calico/cni	v3.1.1	482f47df27e2	3 months ago	68.8MB
us.gcr.io/prp-k8s/perfsonar_testpoint_h	<none>	2fafb9b60974	3 months ago	784MB
us.gcr.io/prp-k8s/perfsonar_testpoint	latest	fec1df3352a3	3 months ago	784MB
quay.io/coreos/kube-rbac-proxy	v0.3.0	543e2018dcac	4 months ago	40.2MB
rook/rook	v0.7.1	ee1353c748c0	4 months ago	435MB
traefik	1.5.3	8c72b944d569	5 months ago	49.6MB
k8s.gcr.io/pause	3.1	da86e6ba6ca1	7 months ago	742kB
quay.io/prometheus/node-exporter	v0.15.2	ff5ecdafc4a2	8 months ago	22.8MB

# Using OpenMPI on K8s: *Using a Custom Container*

```
[paolini@fiona k8s]$ kubectl create -f comet.yaml
deployment.apps "subflow-comet" created

[paolini@fiona k8s]$ ./initcluster.sh sdsu-comet
Generating public/private rsa key pair.
Your identification has been saved in /root/.ssh/id_rsa.
Your public key has been saved in /root/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:cQvHueOgYQI/x0zWRSyEwNyT6lgnNEmcFTNMWZ6pPdM root@subflow-comet-6788dc75d5-778bz
The key's randomart image is:
+---[RSA 2048]-----+
|      *O*Booo      |
|      B.B=.=..     |
|      . +o.B.=     |
|      oo=.o * o     |
|      ++ B S E     |
|      . . = o = .   |
|      . .           |
|      . .           |
|      . .           |
+-----[SHA256]-----+
Enable SSH client HostbasedAuthentication
Enable SSH server HostbasedAuthentication
Enable SSH server rhosts authentication
Enable SSH server HostbasedUsesNameFromPacketOnly
Enable SSH server root login
* Stopping OpenBSD Secure Shell server sshd
* Starting OpenBSD Secure Shell server sshd
.
.
.
# subflow-comet-6788dc75d5-778bz:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-p8z2g:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-pgw78:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-q6wvc:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-778bz:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-p8z2g:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-pgw78:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-q6wvc:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-778bz:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-p8z2g:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-pgw78:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-q6wvc:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-778bz:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-p8z2g:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-pgw78:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
# subflow-comet-6788dc75d5-q6wvc:22 SSH-2.0-OpenSSH_7.2p2 Ubuntu-4ubuntu2.4
```

[ OK ]  
[ OK ]

# Using OpenMPI on K8s: *Initialize Pod cont.*

```
root@subflow-74c57d67d4-bqnc2:/# cat /etc/mpihostfile
```

```
subflow-74c57d67d4-bqnc2 slots=16 max-slots=16
```

```
subflow-74c57d67d4-gd8sn slots=12 max-slots=12
```

```
subflow-74c57d67d4-jjxl5 slots=12 max-slots=12
```

```
subflow-74c57d67d4-jl72q slots=8 max-slots=8
```

```
subflow-74c57d67d4-l9n8t slots=12 max-slots=12
```

```
subflow-74c57d67d4-lj8xw slots=12 max-slots=12
```

```
subflow-74c57d67d4-q8mbj slots=12 max-slots=12
```

```
subflow-74c57d67d4-x5j6r slots=96 max-slots=96
```

**Number of *slots* is the number of processors available to the pod**



# Using OpenMPI on K8s: *MPI Ring Test*

```
[paolini@fiona k8s]$ more RUN
#!/bin/bash

KUBE_NAMESPACE=sdsu
export KUBE_NAMESPACE

MPI_CLUSTER_NAME=subflow
export MPI_CLUSTER_NAME

case "$1" in
  ring)
    kubectl -n $KUBE_NAMESPACE exec -it $MPI_CLUSTER_NAME-74c57d67d4-bqnc2 -- mpirun --allow-run-as-root \
      --hostfile /etc/mpihostfile \
      --mca btl tcp,self \
      -n 8 -npnode 1 --bind-to core \
      /nfs/subflow/k8s/ring
    ;;
```

Direct Open MPI to use TCP-based communications over IP interfaces. Modular Component Architecture (MCA) parameter: byte transfer layer (BTL)

Bind each MPI process to a core.

# Using OpenMPI on K8s: *MPI Ring Test cont.*

```
[paolini@fiona k8s]$ ./RUN ring
process 0 of 8 (on subflow-74c57d67d4-bqnc2)
process 2 of 8 (on subflow-74c57d67d4-jjxl5)
process 1 of 8 (on subflow-74c57d67d4-gd8sn)
process 6 of 8 (on subflow-74c57d67d4-q8mbj)
process 5 of 8 (on subflow-74c57d67d4-lj8xw)
process 3 of 8 (on subflow-74c57d67d4-jl72q)
process 7 of 8 (on subflow-74c57d67d4-x5j6r)
process 4 of 8 (on subflow-74c57d67d4-l9n8t)
token= 8192.0000001, matches TIMES_AROUND*nprocs (things look ok).
doing long send...
...passed long send test.
[paolini@fiona k8s]$
```

Pass a (float) token around the “ring” of 8 processes  $2^{10}$  times. Each process increments the token by +1 when the token is received from its “left” process neighbor, prior to sending the token to its “right” neighbor.

# Benchmarking PRP K8s with LINPACK

- *HPL LINPACK Benchmark*: commonly used performance metric to evaluate a cluster's performance
- Project URL: <https://www.top500.org/project/linpack/>
- FAQs:
  - <http://www.netlib.org/utk/people/JackDongarra/faq-linpack.html>
  - <http://www.netlib.org/benchmark/hpl/faqs.html>
- Download URL: <http://www.netlib.org/benchmark/hpl/>
- Repeatedly solves a linear system of  $N$  equations using LU decomposition with partial row pivoting, giving a performance metric  $R$ .
- We want to identify  $R_{max}$  for a particular  $N$  (with respect to a chosen number of processes and other parameters).

# Benchmarking PRP K8s with LINPACK *cont.*

- Solve a linear system  $Ax = b$  of order  $N$
- Compute LU factorization of  $A$  with row partial pivoting (operations stored in  $P$ ) of the  $n \times (n+1) = N$  coefficient matrix  $[A \ b]$
- $[A \ b] = [[L, U] \ y]$
- $PA=LU \rightarrow LUX=Pb$
- Solve  $Ly=Pb$  for  $y$ , using forward substitution, then solve  $Ux=y$  for  $x$  using backward substitution
- $O(\frac{2}{3}n^3)$  FLOP
- $n \times (n+1)$  coefficient matrix logically partitioned into  $NB \times NB$  blocks, that are cyclically assigned onto a  $P \times Q$  process grid.

# Benchmarking PRP K8s with LINPACK: *Block LU*

- Partition  $A$  into blocks (sub-matrices) of size  $NB \times NB$  and distribute to each process:

$$\begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 & 0 \\ L_{21} & L_{22} & 0 \\ L_{31} & L_{32} & L_{33} \end{bmatrix} \begin{bmatrix} U_{11} & U_{12} & U_{13} \\ 0 & U_{22} & U_{23} \\ 0 & 0 & U_{33} \end{bmatrix}$$

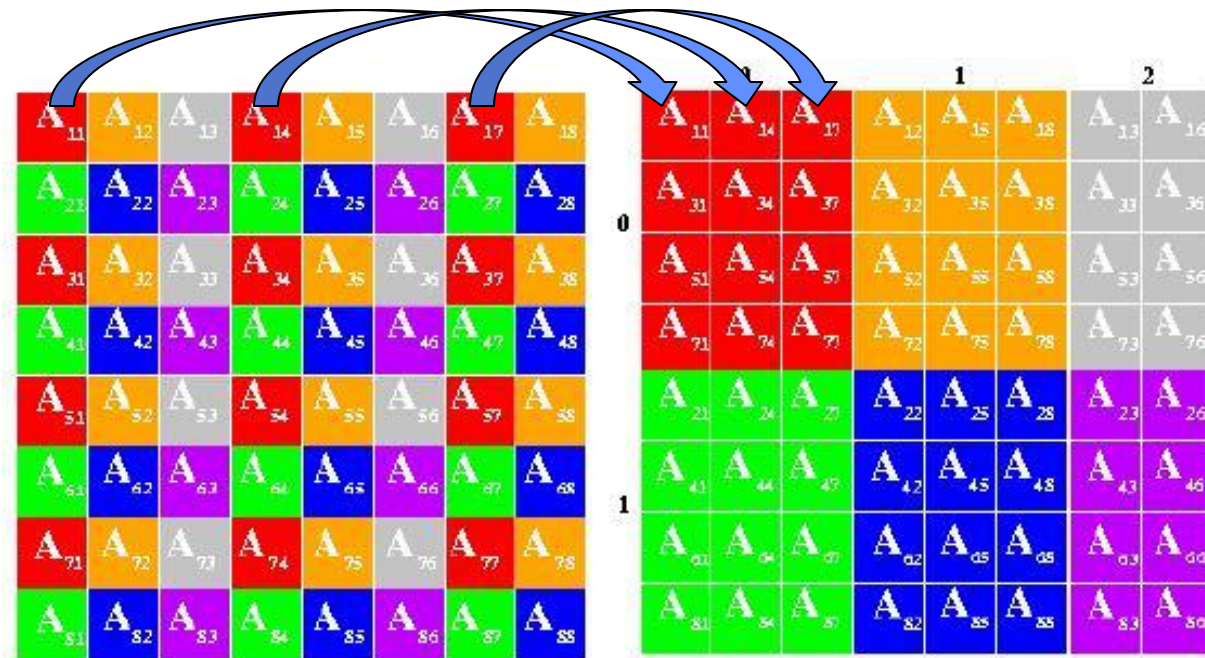
- Partitioning yields the following *block equations* for partitioned LU factorization:

$$\begin{bmatrix} A_{11} = L_{11}U_{11} & A_{12} = L_{11}U_{12} & A_{13} = L_{11}U_{13} \\ A_{21} = L_{21}U_{11} & A_{22} = L_{21}U_{12} + L_{22}U_{22} & A_{23} = L_{21}U_{13} + L_{22}U_{23} \\ A_{31} = L_{31}U_{11} & A_{32} = L_{31}U_{12} + L_{32}U_{22} & A_{33} = L_{31}U_{13} + L_{32}U_{23} + L_{33}U_{33} \end{bmatrix}$$

- Block size,  $NB$ , will affect performance. We can experiment and see how changing the block size yields different performance metrics  $R$  (Gflops) on the PRP K8s cluster.

# Benchmarking PRP K8s with LINPACK: *Processor Grid (PxQ)*

- Example:  $P \times Q = 2 \times 3$ ; each color corresponds to an MPI process assignment.



# Benchmarking PRP K8s with LINPACK: *Launching HPL*

```
[paolini@fiona k8s]$ more RUN
#!/bin/bash
```

```
KUBE_NAMESPACE=sdsu
export KUBE_NAMESPACE
```

```
MPI_CLUSTER_NAME=subflow
export MPI_CLUSTER_NAME
```

```
case "$1" in
```

```
.
```

```
.
```

```
.
```

```
hpl)
```



```
kubectl -n $KUBE_NAMESPACE exec -it $MPI_CLUSTER_NAME-74c57d67d4-bqnc2 -- mpirun --allow-run-as-root \
  --hostfile /etc/mpihostfile \
  --mca btl tcp,self \
  --display-map \
  -n 64 -npnnode 8 --bind-to core:overload-allowed \
  sh -c 'cd /nfs/hpl-2.2/bin/k8s; ./xhpl > xhpl.out 2>&1'
```

```
;;
```



**Max is 8 nodes with 8  
processes per node (64)**

**HPL application is xhpl**

# Benchmarking PRP K8s with LINPACK: *Results*



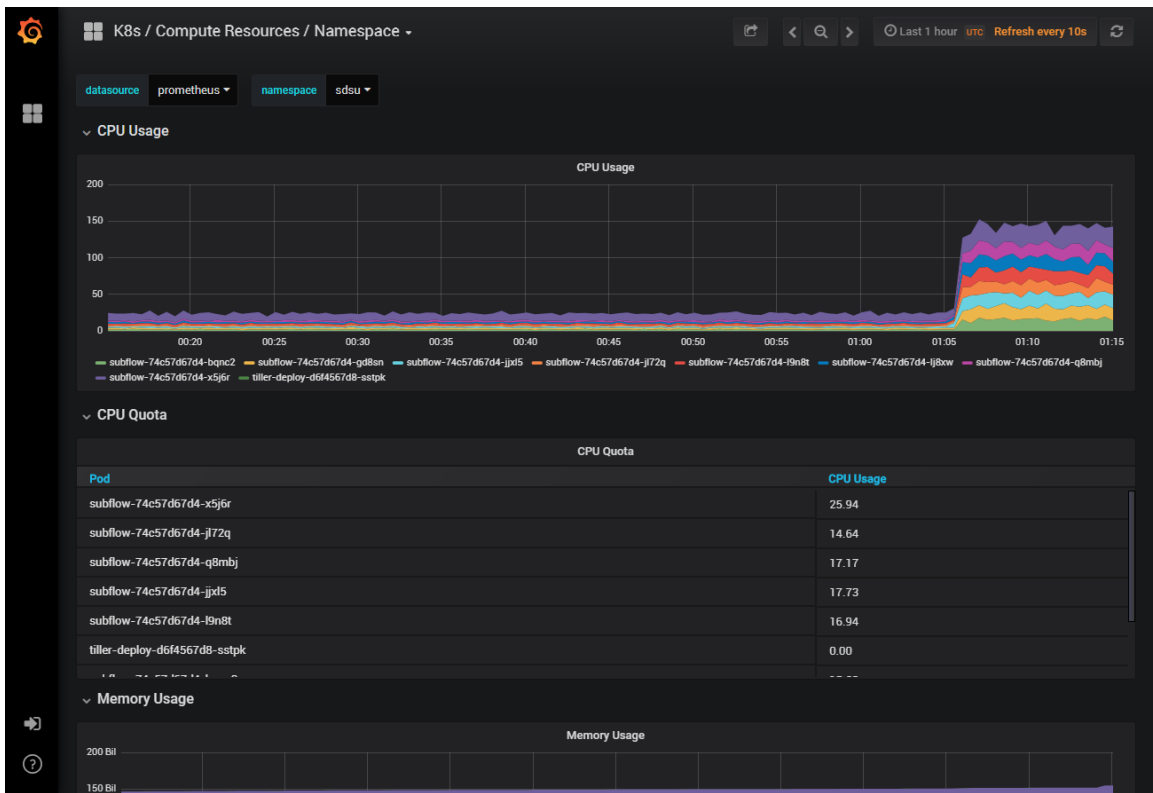
- Blocksize = 256
- (PxQ)=(2x32)
- Panel fact.: Crout
- Recursive Panel fact.: Crout
- 7.056 Gflops





# Benchmarking PRP K8s with LINPACK: *Grafana Monitoring*

- <https://grafana.nautilus.optiputer.net/>
- [https://grafana.nautilus.optiputer.net/d/4r\\_GftHmz/k8s-compute-resources-namespaces?refresh=10s&orgId=1&var-datasource=prometheus&var-namespace=sdsu](https://grafana.nautilus.optiputer.net/d/4r_GftHmz/k8s-compute-resources-namespaces?refresh=10s&orgId=1&var-datasource=prometheus&var-namespace=sdsu)



# Benchmarking PRP K8s with LINPACK: *HPL.dat*

```
root@subflow-74c57d67d4-bqnc2:/nfs/hpl-2.2/bin/k8s# cat HPL.dat
HPLinpac benchmark input file
Innovative Computing Laboratory, University of Tennessee
HPL.out      output file name (if any)
6            device out (6=stdout,7=stderr,file)
9            # of problems sizes (N)
128 256 512 1024 2048 4096 8192 16384 32768 Ns
4            # of NBs
32 64 128 256 NBs
0            PMAP process mapping (0=Row-,1=Column-major)
4            # of process grids (P x Q)
1 2 4 8 Ps
64 32 16 8 Qs
16.0         threshold
1            # of panel fact
1            PFACTs (0=left, 1=Crout, 2=Right)
2            # of recursive stopping criterium
2 4          NBMINs (>= 1)
1            # of panels in recursion
2            NDIVs
1            # of recursive panel fact.
1            RFACTs (0=left, 1=Crout, 2=Right)
1            # of broadcast
0            BCASTs (0=lrg,1=lrM,2=2rg,3=2rM,4=Lng,5=LnM)
1            # of lookahead depth
0            DEPTHs (>=0)
2            SWAP (0=bin-exch,1=long,2=mix)
64           swapping threshold
0            L1 in (0=transposed,1=no-transposed) form
0            U  in (0=transposed,1=no-transposed) form
1            Equilibration (0=no,1=yes)
8            memory alignment in double (> 0)
```

11 values of  $N$ , from  $2^7$  to  $2^{15}$   
such that  $N$  is  $NB$  aligned

Test with 4 block sizes  $NB$ .  
"good" block sizes are almost  
always in the [32 .. 256]  
interval. -- Jack Dongarra


Test with 4 processor grid  
configurations with  $Q \geq P$  and  
 $PQ = 64$ .

Crout's algorithm typically  
performs better than left- and  
right- looking panel methods

# Benchmarking PRP K8s with LINPACK: *Panel Factorization*

```
1          PFACTs (0=left, 1=Crout, 2=Right)
1          # of recursive stopping criterium
4          NBMINs (>= 1)
1          # of panels in recursion
2          NDIVs
1          # of recursive panel fact.
1          RFACTs (0=left, 1=Crout, 2=Right)
```


- Crout matrix decomposition: returns lower triangular matrix and a *unit* upper triangular matrix (Prescott Durand Crout).
- Recursively divides each panel into NDIV subpanels at each step.
- Recursion stops when the current panel is made of less than or equal to NBMIN columns.
- RFACT: method used for recursive sub- panel factorization.
- PFACT: method used after recursive factorization completes.
- Example: NDIV=2, NBMIN=4
- Right- and left-looking factorization algorithms:  
<http://www.netlib.org/utk/papers/outofcore/node3.html>



# Benchmarking PRP K8s with LINPACK: *Right Looking*

- Suppose we have factored the blocks in the first row and column of  $A$ :

$$A = LU = \begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix} = \begin{bmatrix} L_{11} \\ L_{21} \\ L_{31} \end{bmatrix} \begin{bmatrix} U_{11} & U_{12} & U_{13} \\ & U_{22} & U_{23} \\ & & U_{33} \end{bmatrix}$$

 Factored blocks

 Remaining block submatrix to factor

- Since the remaining green block submatrix is independent of the yellow completed region, the green region can be factored in isolation:

$$\begin{bmatrix} A_{22} & A_{23} \\ A_{32} & A_{33} \end{bmatrix} \Rightarrow \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} = P \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} U_{11} & U_{12} \\ 0 & U_{22} \end{bmatrix} = P \begin{bmatrix} L_{11}U_{11} & L_{11}U_{12} \\ L_{21}U_{11} & L_{21}U_{12} + L_{22}U_{22} \end{bmatrix}$$

- Factoring the first column of the block submatrix yields  $L_{11}$ ,  $L_{21}$ , and  $U_{11}$ . We can solve for  $U_{12}$ :

$$A_{12} = L_{11}U_{12} \rightarrow U_{12} = L_{11}^{-1}A_{12}$$


- Finally, we update block  $A_{22}$  and repeat this process:

$$A_{22} = L_{21}U_{12} + L_{22}U_{22} \Rightarrow L_{22}U_{22} = \underbrace{A_{22} - L_{21}U_{12}}_{\text{known}}$$


- Perform LU factorization to solve for  $L_{22}$  and  $U_{22}$

# Benchmarking PRP K8s with LINPACK: *Memory Access*


Left-Looking



Right-Looking




Crout




# Benchmarking PRP K8s with LINPACK: *Comparison*

## Left-Looking




- Blocksize = 256
- (PxQ)=(2x32)
- Panel fact.: Left
- Recursive Panel fact.: Left
- 7.033 Gflops

## Right-Looking



- Blocksize = 256
- (PxQ)=(2x32)
- Panel fact.: Right
- Recursive Panel fact.: Right
- 7.039 Gflops

## Crout



- Blocksize = 256
- (PxQ)=(2x32)
- Panel fact.: Crout
- Recursive Panel fact.: Crout
- 7.056 Gflops

# Benchmarking PRP K8s with LINPACK: SDSC Comet


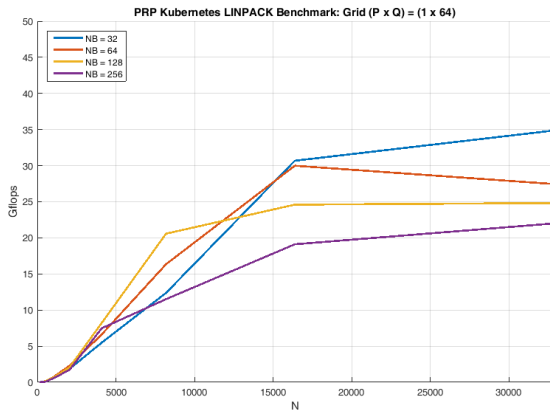
- Comet: eXtreme Science and Engineering Discovery Environment (XSEDE) cluster
- 1,944 nodes, 24 cores/node → 46,656 cores (Haswell)
- FDR InfiniBand Interconnect (MPI latency 1.03-1.97  $\mu$ s)
- K8s on Comet configured by Dmitry Mishin <dmishin@ucsd.edu>

```
[paolini@fiona k8s]$ kubectl create -f comet.yaml
[paolini@fiona k8s]$ kubectl get pods -o wide -n sdsu-comet
```


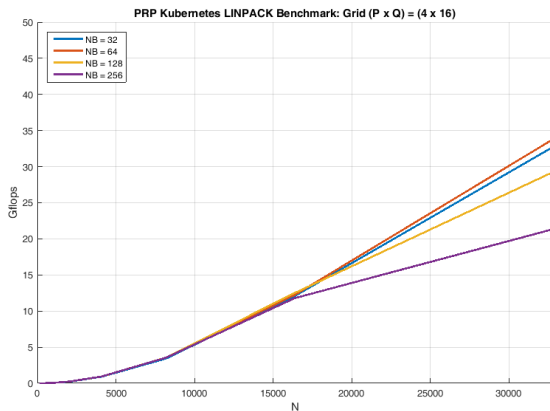
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
subflow-comet-7c8b94f6b6-2hpr6	1/1	Running	0	1h	10.244.66.4	comet-k8s-1.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-b42tb	1/1	Running	0	1h	10.244.68.4	comet-k8s-2.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-gnzgr	1/1	Running	0	1h	10.244.72.4	comet-k8s-3.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-htxxd	1/1	Running	0	1h	10.244.68.3	comet-k8s-2.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-rwc48	1/1	Running	0	1h	10.244.66.3	comet-k8s-1.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-th6qd	1/1	Running	0	1h	10.244.67.5	comet-k8s-0.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-v7cpn	1/1	Running	0	1h	10.244.67.4	comet-k8s-0.sdsc.optiputer.net
subflow-comet-7c8b94f6b6-vfwkh	1/1	Running	0	1h	10.244.72.3	comet-k8s-3.sdsc.optiputer.net



# Benchmarking PRP K8s with LINPACK: Comet Results



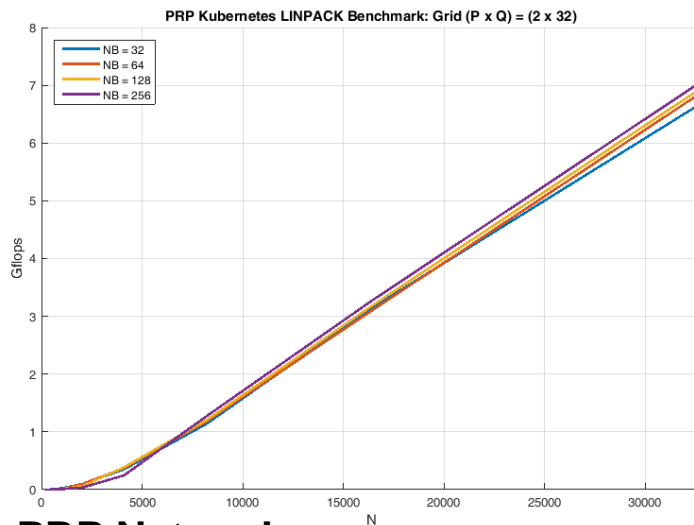
- Blocksize = 32
- (PxQ)=(2x32)
- Panel fact.: Crout
- Recursive Panel fact.: Crout
- 47.54 Gflops






# Benchmarking PRP K8s with LINPACK: *Pros and Cons*

- Slower numerical performance; however:
- Unlimited walltime (so far)
- No formal approval required (so far)
- Required limited administrator intervention



PRP Network

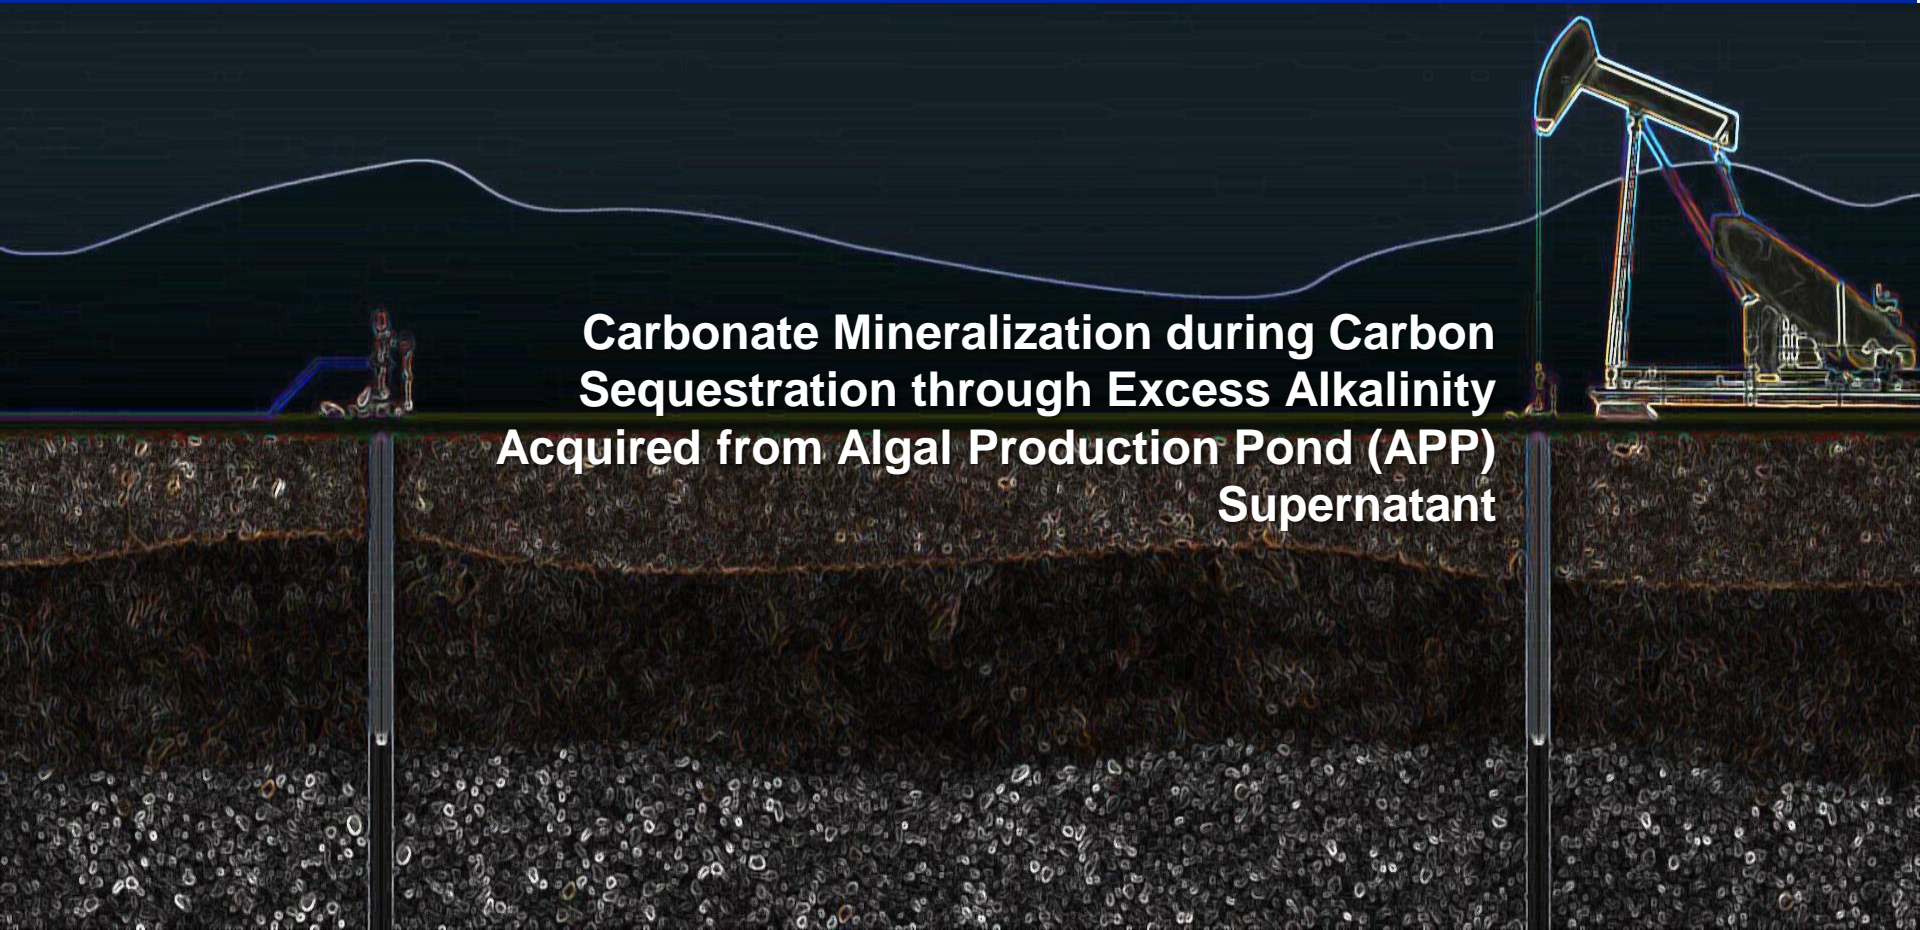
- Superior numerical performance; however:
- Maximum walltime for Comet jobs is 2 days
- Requires XSEDE allocation via approved proposal
- Idle VMs waste allocation SUs
- Required administrator intervention (Dmitry Mishin)



Comet


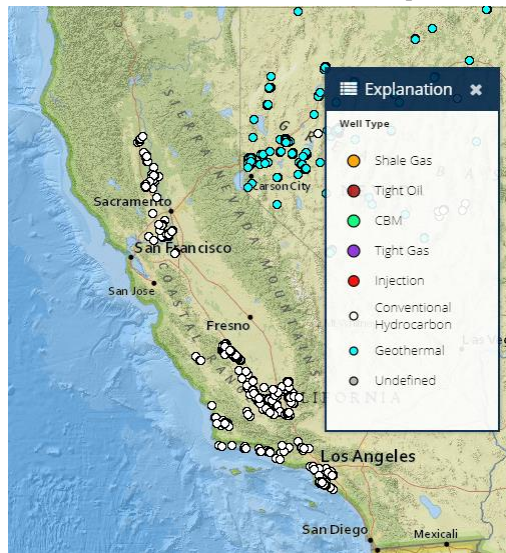
# PRP Science Driver: *Numerical CO<sub>2</sub> Sequestration*

**Carbonate Mineralization during Carbon Sequestration through Excess Alkalinity Acquired from Algal Production Pond (APP) Supernatant**



# Synergistic CCS, EOR, and Biofuel Generation

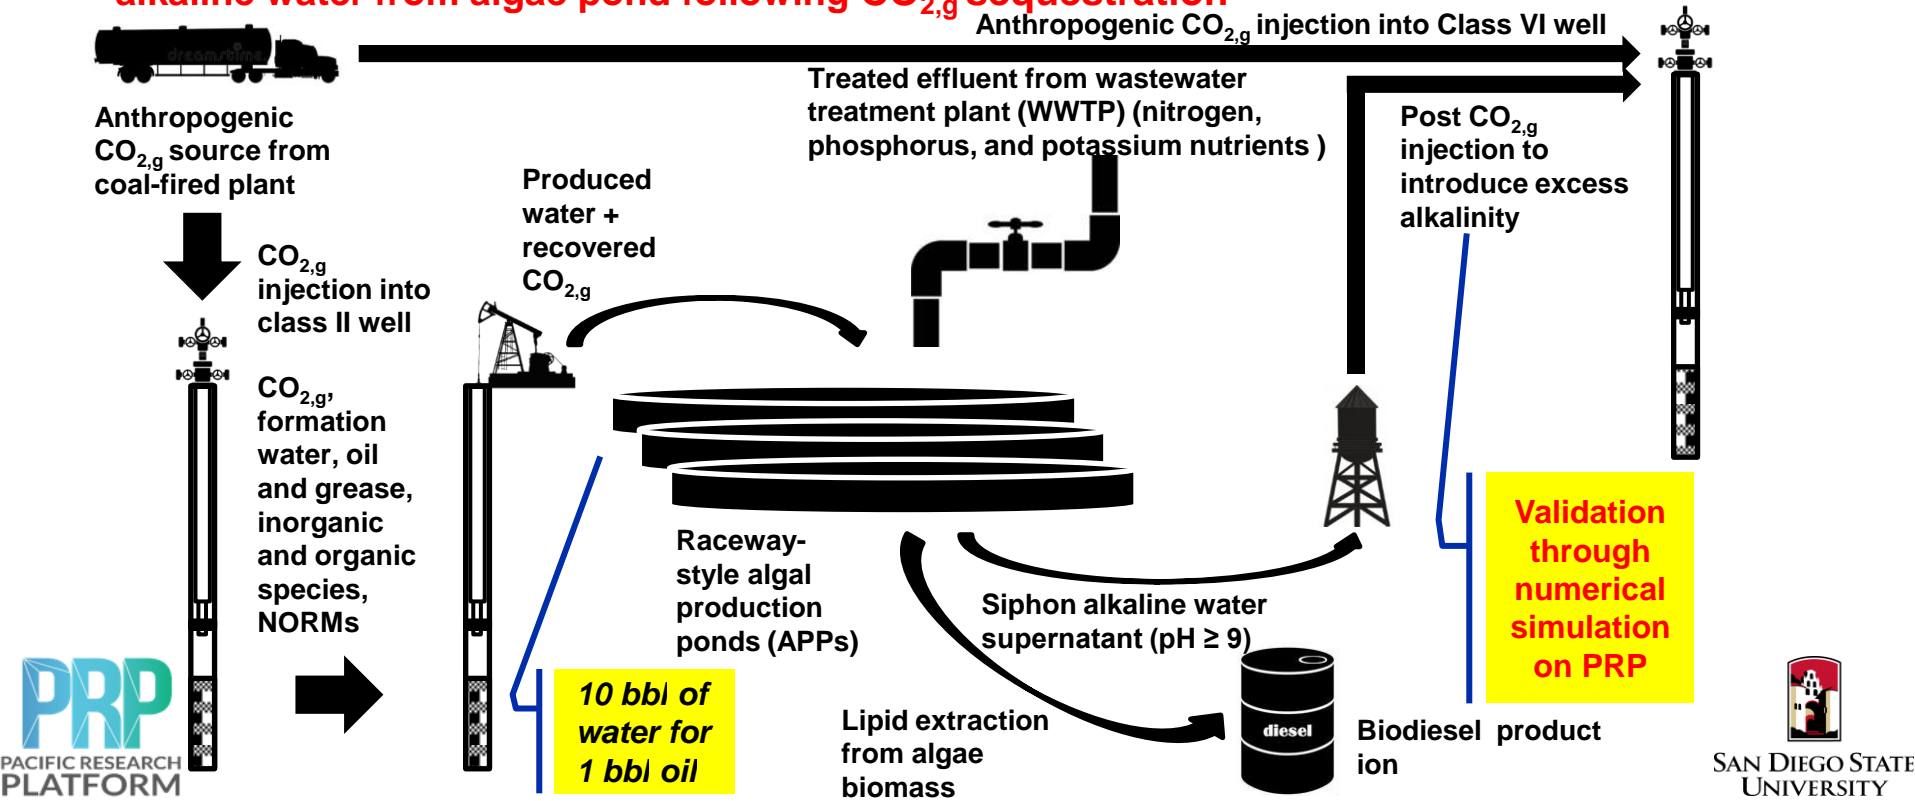
- Exploring Synergies in an Enhanced Waste Management System that Produces Biofuel, Recovers Oil, Sequesters Carbon, and Treats Wastewater



- Primary recovery (natural pressure): 10% original oil in place recovered
- Secondary recovery (water injection): 20% - 40% original oil recovered
- Tertiary EOR (CO<sub>2</sub> injection): 30% to 60%
- US: 114 CO<sub>2</sub> injection projects, 2 billion cubic feet of CO<sub>2</sub> produce 280,000 barrels of oil per day
- Water management: majority (55%) of well operating expenses


# Synergistic CCS, EOR, and Biofuel Generation

- Novel idea:** proposed pilot project to (1) reuse produced water as medium for algal growth to eliminate cost of disposal and (2) **increase rate of carbonate mineralization through injection of alkaline water from algae pond following CO<sub>2,g</sub> sequestration**



# “Locking” Captured CO<sub>2,g</sub> in Place

- Ideal storage mechanism: transform atmospheric CO<sub>2,g</sub> into subsurface solid phase carbonate mineral (e.g. CaCO<sub>3</sub>, MgCO<sub>3</sub>, CaMg(CO<sub>3</sub>)<sub>2</sub>)
- Neutralization of carbonic acid by alkaline earth metals (Ca<sup>2+</sup>, Mg<sup>2+</sup>)
- Solid phase storage: lower risk from upward gas migration and release into atmosphere



- Chemical reaction among dissolved CO<sub>2,g</sub> and mineral solutes in formation water
- “Mineral Trapping”: typically occurs 100 to 10,000 years after injection
- How can we achieve rapid mineralization to lower risk?



# How CO<sub>2</sub> is Sequestered in Rock

- **Structural and Stratigraphic Trapping:**  
CO<sub>2(g)</sub> rises through porous rock until it reaches impermeable cap rock (e.g. shale)
- **Residual-Phase Trapping**  
ScCO<sub>2</sub> stored as residual droplets in (sandstone) rock pore space
- **Dissolution Trapping**  
CO<sub>2</sub> rich water sinks to the bottom of a formation because  $\rho_{\text{H}_2\text{O(l)}+\text{CO}_2\text{(g)}+\text{H}_2\text{CO}_3\text{(aq)}} > \rho_{\text{H}_2\text{O(l)}}$
- **(Ideally) Mineral Trapping**  
Reaction between dissolved CO<sub>2</sub> in formation water with surrounding solutes to create carbonate minerals (e.g. CaCO<sub>3</sub>, MgCO<sub>3</sub>, CaMg(CO<sub>3</sub>)<sub>2</sub>)

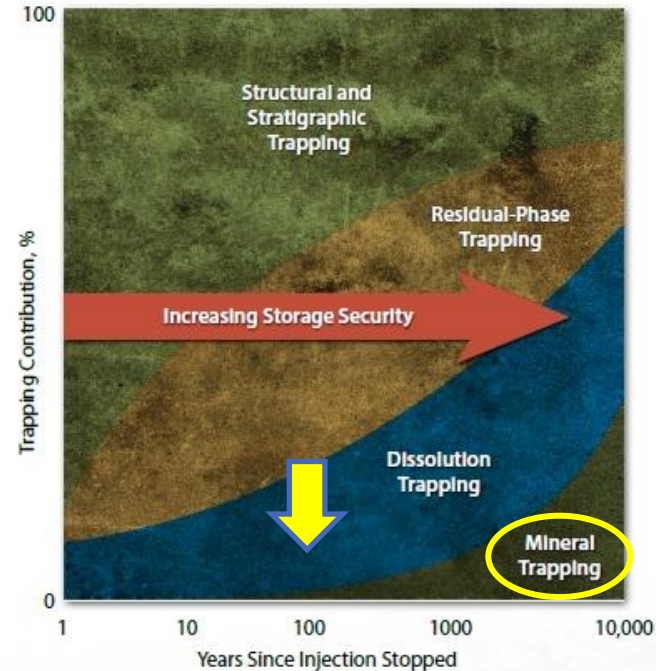






Figure: Plains CO<sub>2</sub> Reduction (PCOR) Partnership Atlas, 4th Edition, Revised, 2013.


Question: how might we **stimulate** a reservoir, post CO<sub>2</sub> injection, to induce carbonate mineral formation **sooner**?

# Problem: Calcite-CO<sub>2</sub> Equilibrium

- CO<sub>2</sub> solvation



- Solid calcite precipitation  $\leftrightarrow$  dissolution



Le Chatelier's principle


- Increase [CO<sub>2,g</sub>] → (1,2) shifts to right → CaCO<sub>3,s</sub> dissolves
- Decrease [CO<sub>2,g</sub>] → (1,2) shifts to left → CaCO<sub>3,s</sub> precipitates
- *Problem:* how to store CO<sub>2,g</sub> and (rapidly) precipitate calcite?

# Carbonate Alkalinity




- **Carbonate alkalinity  $A_C$**

$$A_C = [HCO_3^-] + 2[CO_3^{2-}] + [OH^-] - [H^+]$$

- Alkalinity  $A_C$  increases through carbonate dissolution:



- Carbonate precipitation thus reduces  $A_C$  and increases  $P_{CO_2}$  by converting bicarbonate to  $CO_2$ :



- Thus, carbonate precipitation may not completely sequester carbon. What can be done?



# Excess Alkalinity and Hardness

$$A_c = [HCO_3^-] + 2[CO_3^{2-}] + [OH^-] - [H^+]$$



- Idea 1: introduce external source of alkalinity as  $OH^-$



- Idea 2: introduce external source of hardness as  $Ca^{2+}$  and  $Mg^{2+}$
- Can we show, through modeling and numerical simulation **on a PRP cluster**, increased carbon sequestration as carbonate mineralization from underground injection of excess alkalinity and hardness?

# Excess Alkalinity Source

- Reduction of carbonic acid to formaldehyde and hydroxide production from photosynthesis



High rate algal pond (HRAP) CO<sub>2</sub> addition sump (Craggs et al., 2014)



Seambiotic's commercial algae farm in China uses CO<sub>2</sub> captured from power plant flue gas for photosynthesis

# Produced Water Chemistry

- Produced waters of different geologic ages compared with average concentrations in 35‰ seawater (Collins, A.G. 1975, *Geochemistry of Oilfield Waters*, Elsevier Scientific Publishers, New York. 496 pp)

Source of hardness

Element/Ion	Seawater	Produced Water	
		Highest Concentration (Age <sup>1</sup> )	Range of Mean Concentrations
Salinity	35,000	---	<5000 - >300,000,000
Sodium	10,760	120,000 (J)	23,000 – 57,300
Chloride	19,353	270,000 (P)	46,100 – 141,000
Calcium	416	205,000 (P)	2530 – 25,800
Magnesium	1294	26,000 (D)	530 – 4300
Potassium	387	11,600 (D)	130 – 3100
Sulfate	2712	8400 (T)	210 – 1170
Bromide	87	6000 (J)	46 – 1200
Strontium	0.008	4500 (P)	7 – 1000
Ammonium	---	3300 (P)	23 – 300
Bicarbonate	142	3600 (T)	77 – 560
Iodide	167	1410 (P)	3 – 210
Boron	4.45	450 (T)	8 – 40
Carbonate	---	450 (M)	30 – 450
Lithium	0.17	400 (J)	3 – 50



<sup>1</sup> D, Devonian, J, Jurassic, M, Mississippian, P, Pennsylvanian, T, Tertiary.

(Neff, J. M., et al., Produced Water: Overview of Composition, Fates and Effects, 2014)

# Modeled Reservoir based on Frio Formation

- 0.5 km x 0.5 km x 17.5 m volume
- Three sandstone layers separated by two shale layers
- 4 grid configurations: 20x20x{8, 16, 32, 64}

- Frio location is 30 miles northeast of Houston, in the South Liberty oilfield.
- Injection Zone is a brine-sandstone system with a top seal of 200 feet of Anahuac shale.



# Molar Concentration of $\text{CO}_{2,\text{aq}}$

## Mineralogical Composition

Shale (impervious caprock,  $\kappa=7 \times 10^{-3}$  mD)


Mineral	Volume Fraction
Calcite ( $\text{CaCO}_3$ )	0.20
Quartz ( $\text{SiO}_2$ )	0.28
k-feldspar ( $\text{KAlSi}_3\text{O}_8$ )	0.01
Illite ( $(\text{K},\text{H}_3\text{O})(\text{Al},\text{Mg},\text{Fe})_2(\text{Si},\text{Al})_4\text{O}_{10}[(\text{OH})_2]$ )	0.41

Sandstone (Semi-pervious,  $\kappa= 17.2$  mD)

Mineral	Volume Fraction
Quartz ( $\text{SiO}_2$ )	0.53
k-feldspar ( $\text{KAlSi}_3\text{O}_8$ )	0.11 (subarkose)

10 mM  $\text{CO}_2$ -rich Injectant

Species	Molarity
$\text{H}^+$	$2.6 \times 10^{-6}$
$\text{H}_2\text{O}$	1
$\text{CO}_{2,\text{aq}}/\text{HCO}_3^-/\text{CO}_3^{--}$ complex	0.01



# Modeling Porous Media Flow

- Darcy's Law (fluid flux [g/(cm<sup>2</sup> s)])

$$v_{\text{res}} = -\frac{K}{\mu} \rho_0 (\nabla p_{\text{res}} - \rho g \nabla z)$$

- Fluid Conservation (continuity equation)

$$\partial_t (\rho_0 \phi) + \nabla \cdot v_{\text{res}} = q_{\text{res}}$$

$v$	Fluid mass flux [g/(cm <sup>2</sup> s)]	$g$	Acceleration due to gravity [cm/s <sup>2</sup> ]
$K$	Permeability tensor of porous medium [1/cm]	$\nabla z$	Gravitational direction unit vector [cm]
$\mu$	Dynamic viscosity [g/(cm s)]	$\phi$	Porosity of medium (dimensionless)
$\rho$	Fluid density [g/cm <sup>3</sup> ]	$q$	Fluid source density [g/(cm <sup>3</sup> s)]
$p$	Fluid pressure [gfsc]		

# Poroelastic Models

- **Strain: the symmetric gradient of displacement**

$$\epsilon(u) = \frac{1}{2}(\nabla u + \nabla u^T)$$

- **Rock Stress**

$$\sigma(u) = \lambda(\nabla \cdot u)I + 2G\epsilon(u)$$

- **Poroelastic Stress**

$$\sigma_{por}(u, p) = \sigma(u) - \alpha pI$$

$\lambda$	First Lamé parameter [gfsc]
$G$	Second Lamé parameter (Shear modulus) [gfsc]
$\alpha$	Biot-Willis coefficient (dimensionless)

# Solute Mass Transport Model

- Elemental Conservation of Mass per Unit Volume

$$\frac{\partial e_{\beta}}{\partial t} = \sum_{\alpha=1}^{N_{\alpha}} [\phi D_{\alpha} \nabla^2 c_{\alpha} - \phi \nabla \cdot (c_{\alpha} u)] - \sum_{\gamma=1}^M v_{\beta\gamma} A_{\gamma} G_{\gamma} \quad (\text{Park, 2014})$$

## Elemental mass rate of change term:

rate of increase of concentration of a solute atom  $\beta$  in a fluid element

**Diffusive term:** net rate of increase of solute activity in a fluid element due to diffusive forces

**Advective term:** net rate of flow of solute activity out of a fluid element due to advective forces

**Source term:** net rate of the increase or decrease of a mineral in a fluid element due to chemical kinetics

- Evolution of chemical elemental mass depends on mass-transfer from diffusive and advective forces as well as the precipitation and dissolution of minerals governed by kinetic reaction rates

e elemental mass in fluid [mol/cm<sup>3</sup>]  
 D diffusion coefficient [cm<sup>2</sup>/s]  
 c solute concentration [mol/cm<sup>3</sup>]  
 $\alpha$  aqueous solute species index  
 u fluid velocity [cm/s]

$\phi$  porosity [unitless]  
 $\beta$  solute atom index  
 $\gamma$  mineral index  
 G mineral reaction rate [mol/(cm<sup>2</sup> s)]  
 A mineral spec. surface area [cm<sup>2</sup>/cm<sup>3</sup>]




# Calcite Saturation

Mineral saturation  $\Omega$  is the ratio of the Ion Activity Product ( $IAP$ ) to the (equilibrium) solubility product ( $K_{sp}$ ). For calcite,

Precipitation requires  
 $\Omega > 1$  or  $IAP > K_{sp}$

$$\Omega = \frac{IAP}{K_{sp}} = \frac{a_{Ca^{2+},actual} a_{CO_3^{2-},actual}}{a_{Ca^{2+},equilibrium} a_{CO_3^{2-},equilibrium}}$$




Calcite saturation after 1 day of produced water solution #1

# PRP Results: 1 kg/s, 5 years Equilibrium

Calcite fraction after 1 day of produced water solution #1 and 5 years of equilibrium ( $2\text{e-}7$ )

Solute	Concentration
$\text{OH}^-$	$1\text{e-}2$ M, pH = 12
$\text{Ca}^{++}$	2530 ppm
$\text{Na}^+$	23000 ppm
$\text{Mg}^{++}$	530 ppm
$\text{Cl}^-$	46100 ppm

High alkalinity, low salinity, low hardness test water solution #1




# Comet Results: High Hardness, 1 kg/s, 1 year Equilibrium

Calcite fraction after 1 day of produced water solution #2 and 1 year of equilibrium ( $1.5 \times 10^{-5}$ )

Solute	Concentration
OH <sup>-</sup>	1e-2 M, pH = 12
Ca <sup>++</sup>	25800 ppm
Na <sup>+</sup>	23000 ppm
Mg <sup>++</sup>	4300 ppm
Cl <sup>-</sup>	46100 ppm

High alkalinity, low  
salinity, high hardness  
test water solution #2



High alkalinity, low salinity, high hardness test water solution #2

# PRP Science Driver: *Running Subflow on Comet*

```
demo-comet)
    KUBE_NAMESPACE=sdsu-comet
    export KUBE_NAMESPACE
    HEADNODE=$MPI_CLUSTER_NAME-comet-6788dc75d5-pgw78
    export HEADNODE
    kubect1 -n $KUBE_NAMESPACE exec -it $HEADNODE -- mpirun --allow-run-as-root \
        --hostfile /etc/ssh/shosts.equiv \
        --display-map -n 4 -npnode 1 \
        --mca btl tcp,self \
        sh -c 'cd /nfs/subflow/exe; GRID=FrioTHMC20x20x16; export GRID; ./subflow -omp 4 -i demo.sdb -g $GRID.grid -k $GR
ID > OUTPUT/$GRID-demo.out 2>&1'
    ;;
```

```
[paolini@ps-40g k8s]$ ./RUN.sh demo-comet
```

```
Data for JOB [57040,1] offset 0
```

```
===== JOB MAP =====
```

```
Data for node: subflow-comet-6788dc75d5-778bz Num slots: 24 Max slots: 0 Num procs: 1
Process OMPI jobid: [57040,1] App: 0 Process rank: 0
```

```
Data for node: subflow-comet-6788dc75d5-p8z2g Num slots: 24 Max slots: 0 Num procs: 1
Process OMPI jobid: [57040,1] App: 0 Process rank: 1
```

```
Data for node: subflow-comet-6788dc75d5-pgw78 Num slots: 24 Max slots: 0 Num procs: 1
Process OMPI jobid: [57040,1] App: 0 Process rank: 2
```

```
Data for node: subflow-comet-6788dc75d5-q6wvc Num slots: 24 Max slots: 0 Num procs: 1
Process OMPI jobid: [57040,1] App: 0 Process rank: 3
```

```
=====
```

# Concluding Remarks

Simulations show it may be feasible to induce early carbonate mineral formation through alkaline produced water injection following CO<sub>2</sub> injection

## Further Work

### Intellectual merit

- What species of algae (e.g. *Dunaliella salina*, *Botryococcus braunii*, *Chlorella*, *Dunaliella tertiolecta*, *Gracilaria*, *Pleurochrysis carterae*, *Sargassum*) thrive best, if at all, in produced waters from different oil and gas sites?
- How does produced water salinity affect supernatant alkalinity?
- How do trace amounts of naturally occurring radioactive materials (NORMs) such as Radium affect algae growth? Radium isotopes present in produced water and barite (barium sulfate) scale are <sup>226</sup>Ra and <sup>228</sup>Ra.
- How does alkaline supernatant injection increase the mineralization rates of calcite, magnesite, and dolomite following CO<sub>2,g</sub> injection?
- What is the optimal ratio of produced water to waste water to promote algae growth?

### Broader Impacts

- Cost of produced water disposal is a significant factor in determining the profit of oil and gas production. Reducing or eliminating cost of disposal by diverting produced water to algal production ponds (APPs) for biofuel generation reduces per barrel cost, which lowers fuel cost borne by consumer.
- Transforming atmospheric CO<sub>2,g</sub> to subsurface solid phase carbonate mineral (e.g. CaCO<sub>3</sub>, MgCO<sub>3</sub>, CaMg(CO<sub>3</sub>)<sub>2</sub>) through mineral trapping is a superior technology to mitigate harmful climate-change from fossil fuel combustion. Geologic mineral trapping processes typically occur at the 100yr time scale. Reducing this time reduces risk of unwanted CO<sub>2,g</sub> subsurface transport and leakage back to the atmosphere.

# Questions?